



# IT1313

## Operating Systems and Administration

### Chapter 6

### Process Scheduling

# Instructional Learning Outcomes:

- After this lesson, you should be able to :
  - Explain how scheduling works in the management of processes.
  - Apply the various scheduling strategies and how to evaluate them.

# Scheduling

- Task of managing CPU sharing among a pool of ready processes/threads.
- Possible only with context switching facility.
- The scheduler chooses one of the ready threads to allocate to the CPU when it is available.
- The scheduling policy determines when it is time for a thread to be removed from the CPU and allocate it to another thread.
- The scheduling mechanism determines how the process manager can determine it is time to interrupt the CPU, and how a thread can be allocated to and removed from the CPU.

# Scheduling Mechanisms

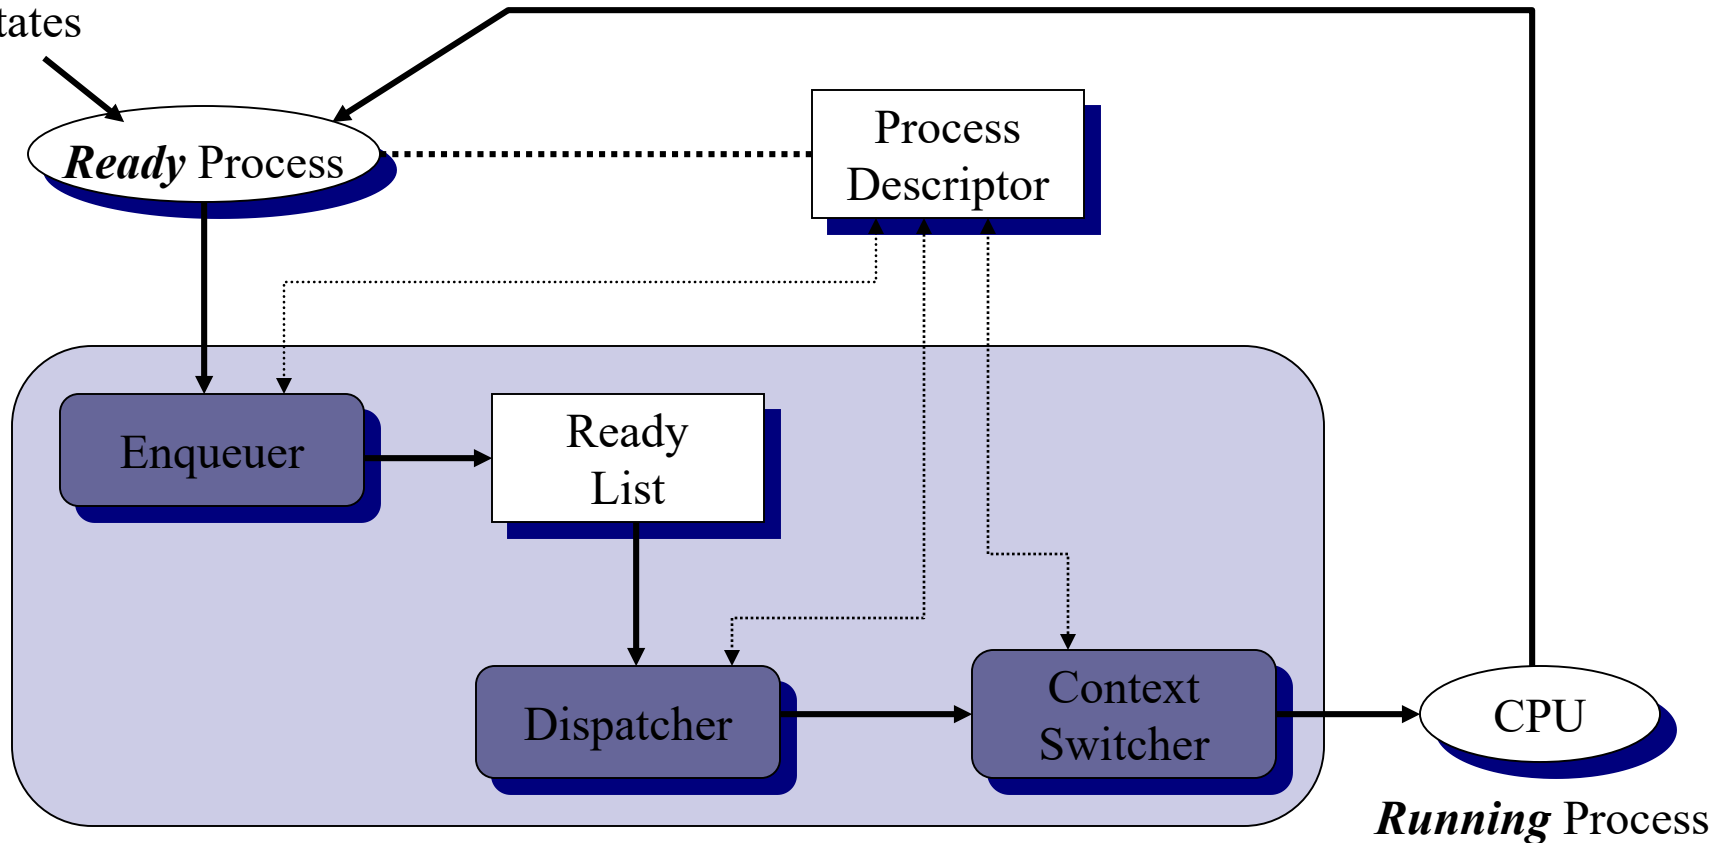
- The scheduling mechanism is composed of three logical parts : the enqueueer, the dispatcher and the context switcher.
  - When a process/thread is changed into the ready state, the **enqueueer** enqueues the corresponding descriptor into a list of processes that are waiting for the CPU (or the ready list). The enqueueer may place the new process anywhere in the list, depending on the **scheduling policy**.

# Scheduling Mechanisms

- When the scheduler switches the CPU from executing one process to another, the **context switcher** saves the contents of all CPU registers (PC, IR, condition status, processor status, and ALU status) for the thread being removed into its descriptor.
- The dispatcher is invoked after the current process is removed from the CPU. The **dispatcher** removes one of the threads from the ready list and then allocates it to the CPU by loading its CPU registers in the thread's descriptor into the CPU.

# The Scheduler

From  
Other  
States



# Scheduling Performance

- Scheduler and its scheduling policy can have a dramatic effect on the performance of a multi-programmed computer.
- The time a process takes to finish execution is dependent on how often it gets to execute on the CPU as despatched by the scheduler.
- Aim is to decrease average time a process takes to execute in the computer and to increase the throughput of a computer in terms of the number of processes get executed per unit time.

# Scheduling Evaluation

- $W(p_i)$  = Total time  $P_i$  spent waiting in ready list (wait time)
- Let  $T_{TRnd}(p_i)$  = Time from  $p_i$  first enter ready to last exit ready (turnaround time)



# Nonpreemptive (voluntary) Schedulers

- First-Come-First-Served (FCFS)
  - Scheduler picks up first job to arrive in the ready list.
- Shortest Job First (SJF) (Nonpreemptive)
  - Scheduler picks up shortest job to arrive in the ready list.
- Priority (PR) (Nonpreemptive)
  - Scheduler picks up the job with highest priority in the ready list.

# Preemptive (involuntary) Schedulers

- Shortest Job First (SJF) (Preemptive)
  - Scheduler picks up shortest job to arrive in the ready list.
- Priority (PR) (Preemptive)
  - Scheduler picks up the job with the highest priority in the ready list.
- Round Robin (RR)
  - Scheduler gives a short time-slice to each job.
- Multi-level Queues
  - Implement multiple queues (eg. Process all foreground processes before background processes, or 80% timeslice to foreground processes, 20% to background processes).

# First-Come First-Served (FCFS)

## ■ Concept

- The process to request first will be allocated the CPU first
- It is non-preemptive
- Using a FCFS queue, PCB of new process will be linked to the tail of the queue
- When CPU is free, process at the head of the FCFS queue will be allocated the CPU

## ■ Advantages:

- Simplest CPU scheduling algorithm

## ■ Disadvantage

- Convoy effect: short process behind long process and waiting for the long process to finished. This results in lower CPU and devices utilization.

# Shortest-Job-First (SJF)

- Concept:
  - Each process is associated with the length of its service time. When the CPU is available, it is assigned to the process that has the smallest (remaining) service time.
- Two methods:
  - nonpreemptive - Once CPU is allocated to a process it cannot be preempted until it completes its service time.
  - preemptive - A current process is preempted if a new process arrives with service time length lesser than the current process's remaining service time.
  - This scheme is also known as the Shortest-Remaining-Time - First (SRTF)
- SJF is optimal – gives minimum average waiting time for a given set of processes.

# Priority Scheduling

## ■ Concept

- A priority number (integer) is associated with each process. The CPU is allocated to the process with the highest priority. Equal- priority processes are scheduled in FCFS order.
  - preemptive
  - nonpreemptive

## ■ Problem: Starvation

- Low priority processes may never be executed.

## ■ Solution: Aging

- Increase the priority of the process as time progresses.

# Round-Robin (RR)

- Each process gets a small unit of CPU time (time quantum), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- New processes are added to the tail of the ready queue, which is a FCFS circular queue.
- Fast response time. It is good for time-sharing system.
- If service time is  $< 1$  time quantum, process release CPU voluntarily.
- If service time  $> 1$  time quantum, context switch will be executed.
- Performance. It depends on size of time quantum:
  - large quantum ( $q$ ). This is equivalent to FIFO.
  - small quantum ( $q$ ). It must be large with respect to the context switch, otherwise overhead is too high.

# Multilevel Queue

- Processes are classified into groups based on some property of the process.
- Ready queue is partitioned into separate queues such as foreground (interactive), background (batch).
- The processes are permanently assigned to one queue.
- Each queue has its own scheduling algorithm.  
For example:
  - RR for foreground queue
  - FCFS for background.

# CPU Schedules


- There is no practical way to determine which policy is the best.
- To study policy performance, we need to consider simplified model, based on a fixed order of processes, which reflect the actual kind of processes a computer can expect.
- To do so, we are provided with :
  - A schedule of processes, their arrival time and CPU service time.
- We need to analyze and produce the following :
  - Gantt chart
  - Average waiting and turnaround time



# CPU Schedules

- A Gantt chart is a 'chart' showing the actual execution of a series of processes by the CPU.
  - It shows the start and end of execution of each process.
  - It does not show the arrival time of each process.
- Waiting time is defined as the total time which a process spends waiting for the CPU in the ready list.
- Turnaround time is the time from the arrival of the process to its completion by the CPU.
  - It includes the time it spent waiting for the CPU.

# CPU Schedules



- The above example shows a typical Gantt chart.
- To find waiting time for p<sub>2</sub>, we take start time for p<sub>2</sub> minus the arrival time of p<sub>2</sub>.
  - Which is  $475 - 0 = 475$ .
- Turnaround time is end time minus the arrival time,
  - which is  $950 - 0 = 950$ .
- Average waiting time is *Total Waiting Time / Number of Processes*
- Average turnaround time is *Total Turanaround Time / Number of Processes*

# CPU Schedules

- In the following slides, we will see how to compute the average turnaround time and average waiting time of the processes for the following scheduling algorithms:
  - FCFS
  - SJF (nonpreemptive)
  - Priority (nonpreemptive)
  - Round Robin
- We assume all the processes arrive at time 0.

# First-Come-First-Served

$i$     $\tau(p_i)$


0   350

1   125


2   475

3   250

4   75




# First-Come-First-Served




# First-Come-First-Served

$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75




# First-Come-First-Served

$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



# First-Come-First-Served

$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75






# FCFS Average Wait Time

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75

- Easy to implement
- Ignores service time, etc
- Not a great performer



$$T_{\text{TRnd}}(p_0) = \tau(p_0) = 350$$

$$W(p_0) = 0$$

$$T_{\text{TRnd}}(p_1) = (\tau(p_1) + T_{\text{TRnd}}(p_0)) = 125 + 350 = 475$$

$$W(p_1) = T_{\text{TRnd}}(p_0) = 350$$

$$T_{\text{TRnd}}(p_2) = (\tau(p_2) + T_{\text{TRnd}}(p_1)) = 475 + 475 = 950$$

$$W(p_2) = T_{\text{TRnd}}(p_1) = 475$$

$$T_{\text{TRnd}}(p_3) = (\tau(p_3) + T_{\text{TRnd}}(p_2)) = 250 + 950 = 1200$$

$$W(p_3) = T_{\text{TRnd}}(p_2) = 950$$

$$T_{\text{TRnd}}(p_4) = (\tau(p_4) + T_{\text{TRnd}}(p_3)) = 75 + 1200 = 1275$$

$$W(p_4) = T_{\text{TRnd}}(p_3) = 1200$$

$$W_{\text{avg}} = (0 + 350 + 475 + 950 + 1200) / 5 = 2974 / 5 = 595$$

$$T_{\text{TRnd}}(\text{avg}) = (350 + 475 + 950 + 1200 + 1275) / 5 = 850$$

# Shortest Job First (Non-preemptive)

$i \quad \tau(p_i)$

0 350

1 125

2 475

3 250


4 75

**Assumption : All jobs arrive at the same time**




# Shortest Job First (Non-preemptive)

$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75




# Shortest Job First (Non-preemptive)

$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75




# Shortest Job First (Non-preemptive)

$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



# Shortest Job First (Non-preemptive)


$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



# Shortest Job First (Non-preemptive)

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75

- Minimizes wait time
- May starve large jobs
- Must know service times



$$T_{\text{TRnd}}(p_0) = \tau(p_0) + \tau(p_3) + \tau(p_1) + \tau(p_4) = 350 + 250 + 125 + 75 = 800$$

$$W(p_0) = 450$$

$$T_{\text{TRnd}}(p_1) = \tau(p_1) + \tau(p_4) = 125 + 75 = 200$$

$$W(p_1) = 75$$

$$T_{\text{TRnd}}(p_2) = \tau(p_2) + \tau(p_0) + \tau(p_3) + \tau(p_1) + \tau(p_4) = 475 + 350 + 250 + 125 + 75 = 1275$$

$$W(p_2) = 800$$

$$T_{\text{TRnd}}(p_3) = \tau(p_3) + \tau(p_1) + \tau(p_4) = 250 + 125 + 75 = 450$$

$$W(p_3) = 200$$

$$T_{\text{TRnd}}(p_4) = \tau(p_4) = 75$$


$$W(p_4) = 0$$

$$W_{\text{avg}} = (450 + 75 + 800 + 200 + 0) / 5 = 1525 / 5 = 305$$

$$T_{\text{TRnd (avg)}} = (800 + 200 + 1275 + 450 + 75) / 5 = 560$$

# Priority Scheduling(Non-preemptive)


$i$	$\tau(p_i)$	Pri
0	350	5
1	125	2
2	475	3
3	250	1
4	75	4






# Priority Scheduling(Non-preemptive)

$i$	$\tau(p_i)$	Pri
0	350	5
1	125	2
2	475	3
3	250	1
4	75	4




# Priority Scheduling(Non-preemptive)

$i$	$\tau(p_i)$	Pri
0	350	5
1	125	2
2	475	3
3	250	1
4	75	4




# Priority Scheduling(Non-preemptive)

i	$\tau(p_i)$	Pri
0	350	5
1	125	2
2	475	3
3	250	1
4	75	4



# Priority Scheduling(Non-preemptive)


i	$\tau(p_i)$	Pri
0	350	5
1	125	2
2	475	3
3	250	1
4	75	4



# Priority Scheduling(Non-preemptive)

i	$\tau(p_i)$	Pri
0	350	5
1	125	2
2	475	3
3	250	1
4	75	4

- Reflects importance of external use
- May cause starvation
- Can address starvation with aging



$$T_{\text{TRnd}}(p_0) = \tau(p_0) + \tau(p_4) + \tau(p_2) + \tau(p_1) + \tau(p_3) = 350 + 75 + 475 + 125 + 250 = 1275$$

$$W(p_0) = 925$$

$$W(p_1) = 250$$

$$W(p_2) = 375$$

$$T_{\text{TRnd}}(p_1) = \tau(p_1) + \tau(p_3) = 125 + 250 = 375$$

$$T_{\text{TRnd}}(p_2) = \tau(p_2) + \tau(p_1) + \tau(p_3) = 475 + 125 + 250 = 850$$

$$T_{\text{TRnd}}(p_3) = \tau(p_3) = 250$$

$$W(p_3) = 0$$

$$T_{\text{TRnd}}(p_4) = \tau(p_4) + \tau(p_2) + \tau(p_1) + \tau(p_3) = 75 + 475 + 125 + 250 = 925$$

$$W(p_4) = 850$$

$$W_{\text{avg}} = (925 + 250 + 375 + 0 + 850) / 5 = 2400 / 5 = 480$$

$$T_{\text{TRnd}}(\text{avg}) = (1275 + 375 + 850 + 250 + 925) / 5 = 735$$

# Round Robin (TQ=50)

$i \quad \tau(p_i)$


0 350

1 125

2 475

3 250

4 75



# Round Robin (TQ=50)

$i \quad \tau(p_i)$


0 350

1 125

2 475

3 250

4 75



# Round Robin (TQ=50)

$i \quad \tau(p_i)$

0 350

1 125

2 475

3 250

4 75

0	100	
$p_0$	$p_1$	$p_2$



# Round Robin (TQ=50)

$i \quad \tau(p_i)$


0 350

1 125

2 475

3 250

4 75



# Round Robin (TQ=50)

$i \quad \tau(p_i)$

0 350

1 125

2 475


3 250

4 75

0	100	200		
$p_0$	$p_1$	$p_2$	$p_3$	$p_4$

# Round Robin (TQ=50)

$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



# Round Robin (TQ=50)

$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75

0	100			200		300		400		475
$p_0$	$p_1$	$p_2$	$p_3$	$p_4$	$p_0$	$p_1$	$p_2$	$p_3$	$p_4$	


# Round Robin (TQ=50)

$i$	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75

0	100	200	300	400	475	550
$p_0$	$p_1$	$p_2$	$p_3$	$p_4$	$p_0$	$p_1$


# Round Robin (TQ=50)

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75




# Round Robin (TQ=50)

i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75



# Round Robin (TQ=50)


i	$\tau(p_i)$
0	350
1	125
2	475
3	250
4	75





# Round Robin (TQ=50)

- |   |             |                                     |
|---|-------------|-------------------------------------|
| i | $\tau(p_i)$ | •Equitable (Fair)                   |
| 0 | 350         | •Most widely-used                   |
| 1 | 125         | •Fits naturally with interval timer |
| 2 | 475         |                                     |
| 3 | 250         |                                     |
| 4 | 75          |                                     |



$$T_{\text{TRnd}}(p_0) = 1100$$

$$T_{\text{TRnd}}(p_1) = 550$$

$$T_{\text{TRnd}}(p_2) = 1275$$

$$T_{\text{TRnd}}(p_3) = 950$$

$$T_{\text{TRnd}}(p_4) = 475$$

$$W(p_0) = 0+200+175+125+100+100+50 = 750$$

$$W(p_1) = 50+200+175 = 425$$

$$W(p_2) = 100+200+150+100+100+100+50 = 800$$

$$W(p_3) = 150+200+150+100+100 = 700$$

$$W(p_4) = 200+200 = 400$$

$$W_{\text{avg}} = (750+425+800+700+400)/5 = 3075/5 = 615$$

$$T_{\text{TRnd}}(\text{avg}) = (1100+550+1275+950+475)/5 = 870$$

# BSD 4.4 (Unix) Scheduling

- Involuntary CPU Sharing
- Preemptive algorithms
- 32 Multi-Level Queues
  - Queues 0-7 are reserved for system functions
  - Queues 8-31 are for user space functions
  - `nice` influences (but does not dictate) queue level

# Windows Scheduling

- Involuntary CPU Sharing across threads
- Preemptive algorithms
- 32 Multi-Level Queues
  - Highest 16 levels are “real-time”
  - Next lower 15 are for system/user threads
    - Range determined by process base priority
  - Lowest level is for the idle thread

# Conclusion

- Scheduling is a requirement for modern OS with multi-programming ability.
- Policy used has a dramatic effect on overall system performance.
- No single 'pure' scheduler is good for all kinds of jobs, so a combination is most often used to provide overall acceptable performance.